

› Web Information Retrieval Boolean Retrieval & Phrase Queries (SOSE 2023)

Frank Hopfgartner, Stefania Zourlidou
Institute for Web Science and Technologies

Credit for these slides

These slides have been adapted from

- Web IR (Zeyd Boukhers-WeST, SOSE 2020)

Recapitulation


- What is Information Retrieval (IR)?
- What makes Web Information Retrieval (WIR) special?
- What are the essential components of IR?
- What is relevance?
- How to evaluate an IR system?
- How to make a corpus for an IR system?
- What is the difference between the various evaluation metrics?

Objectives of this lecture

- What is a *Boolean model*?
 - Examples of Boolean models
- What is a *T-D matrix*?
- What is a *retrieval function*?
- What is an *inverted index*?
- What is a *phrase query*?
- What are *wildcard searches*?

› 1. Boolean Model

Information Retrieval ↔ Fulltext search




Boolean Model

- Boolean
 - Retrieval based on boolean algebra
 - Binary concept of relevance (yes/no)
 - No ranking
 - Queries use boolean operators
- Examples
 - t_1
 - $t_1 \text{ AND } t_2$
 - $t_1 \text{ OR } t_2$
 - $\text{NOT } t_1$
 - $(t_1 \text{ AND } t_2) \text{ OR } (t_3 \text{ AND } t_4)$

Example: documents and queries

- Documents
 - 1. coffee, coffee
 - 2. tea, cup, jar, jar, tea
 - 3. cup, coffee, jar, cup
 - 4. coffee, cup, coffee, jar, tea, cup, coffee, jar, cup, jar
 - 5. jar, water, water, jar
- Queries
 - tea
 - tea AND coffee
 - tea OR coffee
 - NOT coffee
 - (tea AND cup) OR (coffee AND cup)



Theoretic Model

- Corpus: $D = \{d_1, d_2, \dots, d_N\}$
- Vocabulary: $V = \{t_1, t_2, \dots, t_M\}$
- Representation of documents
 - Of interest: is a given term present or not?
 - Document as vector in $\{0,1\}^M$

$$d_i = (\Delta_i^{[1]}, \Delta_i^{[2]}, \dots, \Delta_i^{[M]})^T$$


➤ where

$$\Delta_i^{[j]} = \begin{cases} 0 : t_j \text{ not present} \\ 1 : t_j \text{ present} \end{cases}$$

$$D = \{d_1, d_2, d_3, d_4, d_5\}$$

$$V = \{\text{coffee}, \text{cup}, \text{jar}, \text{tea}, \text{water}\}$$

1. coffee, coffee
2. tea, cup, jar, jar, tea
3. cup, coffee, jar, cup
4. coffee, cup, coffee, jar, tea,
cup, coffee, jar, cup, jar
5. jar, water, water, jar



- Term-document matrix $M \times N$:

$$C = (d_1^{[T]}, d_1^{[T]}, \dots, d_1^{[T]})$$

1. coffee, coffee
2. tea, cup, jar, jar, tea
3. cup, coffee, jar, cup
4. coffee, cup, coffee, jar, tea,
cup, coffee, jar, cup, jar
5. jar, water, water, jar

		d_1	d_2	d_3	d_4	d_5
coffee	t_1	1	0	1	1	0
cup	t_2	0	1	1	1	0
jar	t_3	0	1	1	1	1
tea	t_4	0	1	0	1	0
water	t_5	0	0	0	0	1

- Set of all possible queries: Q
- Recursive definition of queries
 - Each term is a query
 - $\forall t \in V : t \in Q$
 - Combination of queries are queries
 - $\forall q_1, q_2 \in Q : q_1 \text{ AND } q_2 \in Q$
 - $\forall q_1, q_2 \in Q : q_1 \text{ OR } q_2 \in Q$
 - $\forall q \in Q : \text{NOT } q \in Q$
 - For clarity, add parenthesis
 - $\forall q \in Q : (q) \in Q$

Retrieval function ρ

- Function

$$\rho : D \times Q \rightarrow \{0,1\}$$

- Values:

- $\rho(d, q) = 1$: document d is relevant to query q
 - $\rho(d, q) = 0$: document d is not relevant to query q

- Recursive definition

- For single term queries

- $\rho(d_i, t_j) = \Delta_i^{[j]}$

- For compositions

- $\rho(d_i, q_1 \text{AND} q_2) = \min(\rho(d_i, q_1), \rho(d_i, q_2))$

- $\rho(d_i, q_1 \text{OR} q_2) = \max(\rho(d_i, q_1), \rho(d_i, q_2))$

- $\rho(d_i, \text{NOT} q) = 1 - \rho(d_i, q)$

Alternative View using T-D Matrix

- Single term query
 - Result: row in T-D-Matrix
- Combination
 - Bit operations on rows

		d_1	d_2	d_3	d_4	d_5
coffee	t_1	1	0	1	1	0
cup	t_2	0	1	1	1	0
jar	t_3	0	1	1	1	1
tea	t_4	0	1	0	1	0
water	t_5	0	0	0	0	1

- Example:
 - coffee AND tea

This view will be
exploited for realization

		d_1	d_2	d_3	d_4	d_5
coffee	t_1	1	0	1	1	0
tea	t_4	0	1	0	1	0
AND						
RESULT		0	0	0	1	0


› 2. Inverted Index & Boolean Retrieval

Inverted Index


- Central data-structure in IR
- Requirements from T-D-matrix view
 - Lookup row vectors (term-vector)
 - Apply bit operations (AND, OR, NOT)
- Additionally
 - T-D-Matrix is typically very large
 - 1,000,000 documents
 - 100,000 terms
 - each entry 1 bit
 - entire Matrix: 12.5 GB
 - T-D-Matrix is typically very sparse
 - 1,000 terms per document: 99% of entries are 0
 - Compress matrix : store only entries with value 1

Inverted Index

- Data structure consisting of
 - Lookup terms (row vector)
 - Search tree




- Posting-List of non-zero entries in vector
 - Linked list of postings



- Posting: reference to a document

Inverted index construction




Document icons from free icon set: <http://www.icojoy.com/articles/44/>

- Tokenization
 - Cut character sequence into word tokens
 - Deal with "*John's*", **a state-of-the-art solution**
- Normalization
 - Map text and query term to same form
 - You want **U.S.A.** and **USA** to match
- Stemming
 - We may wish different forms of a root to match
 - **authorize, authorization**
- Stop words
 - We may omit very common words (or not)
 - **the, a, to, of**


Inverted Index Construction

- Sort terms with document references




Inverted Index

- Build search tree over vocabulary
- Compile a posting list for each term



Implementing Boolean Retrieval




- Search for single term $q = t$
 - Look up posting list for t ... and done!
- Example: query „tea“


Result of a query is a
posting list



Implementing Boolean Retrieval




- Search for $q = q_1 \text{ AND } q_2$
 - Intersect the result lists (posting lists)
- Example: query „coffee AND tea“



Implementing Boolean Retrieval




- Search for $q = q_1 \text{ OR } q_2$
 - Merge the result lists (posting lists)
- Example: query „coffee OR tea“




Implementing Boolean Retrieval

- Query: NOT q ?
 - Technically simple
 - Invert posting list
- In practice
 - Too many entries
- BUT operator
 - Binary operator
 - Defined as: $q_1 \text{ BUT } q_2 = q_1 \text{ AND } (\text{NOT } q_2)$
- Example: „tea BUT coffee“




Implementing Boolean Retrieval

- More complex queries:
 - $t_1 \text{ AND } t_2 \text{ OR } t_3 \text{ AND } t_4$
 - Recursive approach




Optimization

- Intersect short lists first
 - Faster to process
 - Result lists get shorter
 - Empty list serves as a stop criterion
- Example: query „coffee AND jar AND water“



- Annotate terms with the length of posting list
 - List of length = number of documents containing term
 - Document frequency



› 3. Phrase Queries

- We want to be able to answer queries such as
"stanford university" – as a phrase
- Thus the sentence ***"I went to university at Stanford"*** is not a match
 - The concept of phrase queries has proven easily understood by users; one of the few “advanced search” ideas that works
 - Many more queries are *implicit phrase queries*
- For this, it no longer sufficient to store only
<term : docs> entries

A first attempt: Biword indexes

- Index every consecutive pair of terms in the text as a phrase
- For example, the text “Friends, Romans, Countrymen” would generate the biwords
 - *friends romans*
 - *romans countrymen*
- Each of these biwords is now a dictionary term
- Two-word phrase query-processing is now direct

- Longer phrases can be processed by breaking them down
- ***stanford university palo alto*** can be broken into the Boolean query on biwords:
 - ***stanford university AND university palo AND palo alto***
- Without the docs, we cannot verify that the docs matching the above Boolean query do contain the phrase
- It can have false positives!

- False positives, as noted before
- Index blow-up due to bigger dictionary
 - Infeasible for more than biwords, big even for them
- Biword indexes are not the standard solution (for all biwords), but can be part of a compound strategy

Solution 2: Positional indexes

- In the postings, store, for each **term**, the position(s) in which tokens of it appear

<**term**, number of docs containing **term**:

doc1: position1, position2 ... ;

doc2: position1, position2 ... ;

etc.>

Example: positional index

<**be**: 993427;
1: 7, 18, 33, 72, 86, 231;
2: 3, 149;
4: 17, 191, 291, 430, 434;
5: 363, 367, ...>


Which of docs **1, 2, 4, 5**
could contain "**to be**
or not to be"?

- For phrase queries, we use a merge algorithm recursively at the document level

- Extract inverted index entries for each distinct term:
to, be, or, not
- Merge their *doc: position* lists to enumerate all positions with "***to be or not to be***".
 - ***to***
 - 2:1,17,74,222,551; **4:8,16,190,429,433**; 7:13,23,191; ...
 - ***be***
 - 1:17,19; **4:17,191,291,430,434**; 5:14,19,101; ...
- Same general method for proximity searches

Position Index

- Store position of term in postings



- Intersect position-lists with offset
- Also allows for NEAR operator

- LIMIT! /3 STATUTE /3 FEDERAL /2 TORT
 - Here, $/k$ means “within k words of”
- Clearly, positional indexes can be used for such queries;
 - biword indexes cannot


- A positional index expands postings storage *substantially*
 - Even though indices can be compressed
- Nevertheless, a positional index is now standardly used because of the power and usefulness of phrase and proximity queries
 - ... whether used explicitly or implicitly in a ranking retrieval system

➤ 4. Wildcard Searches

- Flexible search
 - Leave out fragments of terms, mark with wildcard
 - Question mark (?): single character
 - Asterisk (*): zero, one or more characters
- Examples:
 - na?ve → naive, naïve
 - universit* → university, universität, universitá
 - go* → go, goes, gone
 - g?n* → gun, gone, gin, ginger

Wildcard

- Simple case
 - Use sorting in inverted index
 - Example: „col*“
- Further application: Autocomplete (without ranking)



Range selection

ginger

ginger

gin

ginkgo

ginseng

gingrich

About 167,000,000 results (0.17 seconds)

Advanced search

› 5. Summary

Summary

- At the end of this lecture, you are expected to understand the concepts of
 - Boolean model
 - T-D matrix
 - Retrieval function
 - Inverted index
 - Posting list
 - Optimization
 - Phrase queries
 - Positional index
 - Wildcard search

References

- [1] <https://olat.vcrp.de/url/RepositoryEntry/2565867182>
- [2] <https://videoakademie.ko-id.de/Panopto/Pages/Sessions>List.aspx?folderID=07fcee3e-4b21-482c-90ab-ab9500ec2019>
- [3] <http://west.uni-koblenz.de/studying/ws1920/machine-learning-and-data-mining>
- [4] <https://videoakademie.ko-id.de/Panopto/Pages/Viewer.aspx?id=545f21ba-a671-4ada-b137-ab9500f21941>
- [5] <https://nlp.stanford.edu/IR-book/information-retrieval-book.html> Christopher D. Manning, Prabhakar Raghavan and Hinrich Schütze, Introduction to Information Retrieval, Cambridge University Press. 2008.
- [6] <https://www.gartner.com/>