

➤ Web Information Retrieval

Vector Space Model & Relevance Feedback

(SOSE 2023, 10.05.2023)

Frank Hopfgartner, Stefania Zourlidou
Institute for Web Science and Technologies

Credit for these slides

These slides have been adapted from

- Web IR (Zeyd Boukhers-WeST, SOSE 2020)

Recapitulation

- What is a Boolean model?
- What is inverted index?
- What are phrase queries?
- What are proximity queries?

Objectives of this lecture

- Boolean Model: Pros and Cons
- Ranked retrieval model
- Documents scoring
 - TF-IDF
- Query-document matching
 - Jaccard
 - Cosine
- Vector Space Model
- Relevance feedback

› 1. Boolean vs. Ranked Retrieval Models

Boolean Model: pros and cons

- Advantages
 - Good for expert users with precise understanding of their needs and the collection
 - Good for applications which can easily consume (process) 1000s of results
- Disadvantages
 - Not good for the majority of users: expressing information needs as Boolean expressions is unintuitive
 - Not practical for users: no ranking + too many results

- Boolean queries often result in either too few (=0) or too many (1000s) results
- Query 1: “*standard user dlink 650*” → 200,000 hits
- Query 2: “*standard user dlink 650 no card found*”: 0 hits
- It takes a lot of skill to come up with a query that produces a manageable number of hits
 - AND gives too few; OR gives too many

- Rather than a set of documents satisfying a query expression, in ranked retrieval, the system returns an ordering over the (top) documents in the collection for a query
- Free text queries: Rather than a query language of operators and expressions, the user's query is just one or more words in a human language
- In principle, there are two separate choices here, but in practice, ranked retrieval has normally been associated with free text queries and vice versa

- When a system produces a ranked result set, large result sets are not an issue
 - Indeed, the size of the result set is not an issue
 - We just show the top k (≈ 10) results
 - We don't overwhelm the user
- Premise: the ranking algorithm works

Scoring as the basis of ranked retrieval

- We wish to return in order the documents most likely to be useful to the searcher
 - How can we rank-order the documents in the collection with respect to a query?
 - Assign a score – say in $[0, 1]$ – to each document
 - This score measures how well the document and query “match”



- We need a way of assigning a score to a query-document pair
 - Let's start with a one-term query
 - If the query term does not occur in the document: score should be 0
 - The more frequent the query term in the document is, the higher the score (should be)
 - We will consider a number of alternatives to this

Jaccard coefficient

- A commonly used measure of overlap of two sets
- Let A and B be two sets
 - Jaccard coefficient

$$\text{JACCARD}(A, B) = \frac{|A \cap B|}{|A \cup B|}$$

- $\text{JACCARD}(A, A) = 1$
- $\text{JACCARD}(A, B) = 0$ if $A \cap B = 0$
- A and B don't have to be the same size
- Always assigns a number between 0 and 1

Jaccard coefficient: Example

- What is the query-document match score that the Jaccard coefficient computes for the
 - Query: “ides of March”, and
 - Document “Caesar died in March”
 - $\text{JACCARD}(q, d) = 1/6$

Issues with Jaccard for scoring



$$\text{JACCARD}(A, B) = \frac{|A \cap B|}{|A \cup B|}$$

- It doesn't consider *term frequency* (how many times a term occurs in a document)
- Are all words in a document equally important? (*stop words*)
 - Rare terms in a collection are more informative than frequent terms
 - Jaccard does not consider this aspect
- We need a more sophisticated way of normalizing for length

Binary term-document incidence matrix

	Antony and Cleopatra	Julius Caesar	The Tempest	Hamlet	Othello	Macbeth
Antony	1	1	0	0	0	1
Brutus	1	1	0	1	0	0
Caesar	1	1	0	1	1	1
Calpurnia	0	1	0	0	0	0
Cleopatra	1	0	0	0	0	0
mercy	1	0	1	1	1	1
worser	1	0	1	1	1	0

- Each document is represented by a binary vector $\in \{0,1\}^{|V|}$

Term-document count matrices

- Consider the number of occurrences of a term in a document
 - Each document is a count vector in \mathbb{N}^v : a column below

	Antony and Cleopatra	Julius Caesar	The Tempest	Hamlet	Othello	Macbeth
Antony	157	73	0	0	0	0
Brutus	4	157	0	1	0	0
Caesar	232	227	0	2	1	1
Calpurnia	0	10	0	0	0	0
Cleopatra	57	0	0	0	0	0
mercy	2	0	3	5	5	1
worser	2	0	1	1	1	0

Bag of words model

- Vector representation doesn't consider the ordering of words in a document
- *John is quicker than Mary* and *Mary is quicker than John* have the same vectors
- This is called the bag of words model
- In a sense, this is a step back: The positional index was able to distinguish these two documents
- We will look at “recovering” positional information later in this course
- For now: bag of words model

Term frequency tf

- The term frequency $\text{tf}_{t,d}$ of term t in document d is defined as the number of times that t occurs in d
- We want to use **tf** when computing query-document match scores
But how?
- Raw term frequency is not what we want
 - A document with 10 occurrences of the term is more relevant than a document with 1 occurrence of the term
 - But not 10 times more relevant
- **Relevance does not increase proportionally with term frequency**

- The log frequency weight of term t in d is defined as follows

$$w_{t,d} = \begin{cases} 1 + \log_{10} tf_{t,d} & \text{if } tf_{t,d} > 0 \\ 0 & \text{Otherwise} \end{cases}$$

- $tf_{t,d} \rightarrow w_{t,d}$
 - 0 → 0, 1 → 1, 2 → 1.3, 10 → 2, 1000 → 4, etc.
- Score for a document-query pair: sum over terms t in both q and d:
 - tf-matching-score(q, d) = $\sum_{t \in q \cap d} (1 + \log tf_{t,d})$
- The score is 0 if none of the query terms is present in the document

- Rare terms are more informative than frequent terms
 - Recall stop words
- Consider a term in the query that is rare in the collection (e.g., *archaeology*)
- A document containing this term is very likely to be relevant to the query *archaeology*
 - We want a high weight for rare terms like *archaeology*

Document frequency, continued

- Frequent terms are less informative than rare terms
- Consider a query term that is frequent in the collection (e.g., *high*, *increase*, *line*)
- A document containing such a term is more likely to be relevant than a document that doesn't
- But it's not a sure indicator of relevance
- For frequent terms, we want high positive weights for words like *high*, *increase*, and *line*
- But lower weights than for rare terms
- We will use document frequency (df) to capture this

Inverse document frequency (idf) weight

- df_t is the document frequency of term t : the number of documents in the collection that contain a term t
 - df_t is an inverse measure of the informativeness of t
 - $\text{df}_t \leq N$
- We define the **idf** (inverse document frequency) of t by

$$\text{idf}_t = \log (N/\text{df}_t)$$

- We use $\log (N/\text{df}_t)$ instead of N/df_t to “dampen” the effect of idf

idf example, suppose $N = 1$ million

term	df_t	idf_t
calpurnia	1	6
animal	100	4
sunday	1,000	3
fly	10,000	2
under	100,000	1
the	1,000,000	0

$$idf_t = \log_{10} (N/df_t)$$

- There is one idf value for each term t in a collection

Effect of idf on ranking

- Does idf have an effect on ranking for one-term queries, like
 - iPhone
- idf has no effect on ranking one term queries
 - idf affects the ranking of documents for queries with at least two terms
 - For the query *capricious person*, *idf* weighing makes occurrences of *capricious* count for much more in the final document ranking than occurrences of *person*

Collection vs. Document frequency

- The collection frequency of t is the number of occurrences of t in the collection, counting multiple occurrences
- Example

Word	Collection frequency	Document frequency
Insurance	10440	3997
Try	10422	8760

- Which word is a better search term (and should get a higher weight)?

- The **tf-idf_{t,d}** weight of a term t in a document d is the product of its tf weight and its idf weight

$$\begin{aligned} \text{tf-idf}_{t,d} &= \text{tf}_{t,d} \times \text{idf}_{t,d} \\ &= \text{tf}_{t,d} \times \log\left(\frac{N}{\text{df}_t}\right) \end{aligned}$$

- Best known weighting scheme in information retrieval
 - Note: the “-” in tf-idf is a hyphen, not a minus sign!
 - Alternative names: tf.idf, tf x idf
- Increases with the number of occurrences within a document
- Increases with the rarity of the term in the collection

Score for a document given a query

$$Score(q, d) = \sum_{t \in q \cap d} \text{tf. idf}_{t,d}$$

- There are many variants based on
 - how “tf” is computed (with/without logs)
 - whether the terms in the query are weighted too
 - and more

› 2. Vector Space Model

Binary term-document incidence matrix

	Antony and Cleopatra	Julius Caesar	The Tempest	Hamlet	Othello	Macbeth
Antony	1	1	0	0	0	0
Brutus	1	1	0	1	0	0
Caesar	1	1	0	1	1	1
Calpurnia	0	1	0	0	0	0
Cleopatra	1	0	0	0	0	0
mercy	1	0	1	1	1	1
worser	1	0	1	1	1	0

Each document is represented by a binary vector $\in \{0,1\}^{|V|}$

Count matrices

	Antony and Cleopatra	Julius Caesar	The Tempest	Hamlet	Othello	Macbeth
Antony	157	73	0	0	0	0
Brutus	4	157	0	1	0	0
Caesar	232	227	0	2	1	1
Calpurnia	0	10	0	0	0	0
Cleopatra	57	0	0	0	0	0
mercy	2	0	3	5	5	1
worser	2	0	1	1	1	0

Binary → count → weight matrix

	Antony and Cleopatra	Julius Caesar	The Tempest	Hamlet	Othello	Macbeth
Antony	5,25	3,18	0	0	0	0
Brutus	1,21	6,1	0	1	0	0
Caesar	8,59	2,54	0	1,51	0,25	1,95
Calpurnia	0	1,54	0	0	0	0
Cleopatra	2,85	0	0	0	0	0
mercy	1,51	0	1,9	0,12	5,25	0,88
worser	1,37	0	0,11	4,15	0,25	0


- Each document is now represented by a real-valued vector of tf-idf weights $\in \mathbb{R}^{|V|}$

- So we have a $|V|$ -dimensional vector space
- Terms are axes of the space
- Documents are points or vectors in this space
- Very high-dimensional: tens of millions of dimensions when you apply this to a web search engine
- These are very sparse vectors - most entries are zero

Idea



- Documents as vectors in a high dimensional space
 - Queries are documents
- Terms define dimensions (base vectors)
- Assumptions
 - Base vectors orthogonal
 - Similar documents have similar vectors
 - Vector similarity indicates document similarity
 - Distance
 - Angle



Similarity ranking for q ?


- With distance
- With angle

Formalizing vector space proximity

- First cut: distance between two points
 - (= distance between the end points of the two vectors)
- Euclidean distance?
 - Euclidean distance is a bad idea . . .
 - ... because Euclidean distance is large for vectors of different lengths

Why distance is a bad idea

- The Euclidean distance between \vec{q} and \vec{d}_2 is large even though the distribution of terms in the query \vec{q} and the distribution of terms in the document \vec{d}_2 are very similar




Use angle instead of distance

- Thought experiment: take a document d and append it to itself. Call this document d'
- “Semantically” d and d' have the same content
- The Euclidean distance between the two documents can be quite large
- The angle between the two documents is 0, corresponding to maximal similarity
- Key idea: Rank documents according to angle with query

From angles to cosines

- Cosine is a monotonically decreasing function for the interval $[0^\circ, 180^\circ]$



The following two notions are equivalent

- Rank documents in decreasing order of the angle between query and document
- Rank documents in increasing order of cosine(query,document)

- But how – *and why* – should we compute cosines?

cosine(query, document)

$$\cos(\vec{q}, \vec{d}) = \frac{\vec{q} \bullet \vec{d}}{\|\vec{q}\| \|\vec{d}\|} = \frac{\vec{q}}{\|\vec{q}\|} \bullet \frac{\vec{d}}{\|\vec{d}\|} = \frac{\sum_{i=1}^{|V|} q_i d_i}{\sqrt{\sum_{i=1}^{|V|} q_i^2} \sqrt{\sum_{i=1}^{|V|} d_i^2}}$$

Dot product Unit vectors

- q_i is the tf-idf weight of term i in the query
- d_i is the tf-idf weight of term i in the document

$\cos(\vec{q}, \vec{d})$ is the cosine similarity of q and d ... or,
equivalently, the cosine of the angle between q and d

Vector Space Model (VSM) Framework

- Corpus: $D = \{d_1, \dots, d_N\}$
- Vocabulary: $V = \{t_1, \dots, t_M\}$
- Documents as vectors in R^M

$$\vec{d}_i = (w_i^{[1]}, w_i^{[2]}, \dots, w_i^{[M]})$$

- Where $w_i^{[j]}$ weight of term t_j in document d_i
- Queries:

$$\vec{q} = (w_q^{[1]}, w_q^{[2]}, \dots, w_q^{[M]})$$

- Retrieval function

$$\rho(\vec{d}_i, \vec{q}) = \text{sim}(\vec{d}_i, \vec{q})$$

- Similarity measure

$$\text{sim} : \mathbf{R}^M \times \mathbf{R}^M \rightarrow [0, 1]$$

- With
 - Value 1: same vector
 - Value 0: „completely different“ vector

- Result list

- All documents with $\rho > 0$, sorted by descending score

Combined Weights

- Combined weight: TF-IDF

$$\begin{aligned} w_{\text{TF.IDF}}(t_j, d_i) &= w_{\text{local}}(t_j, d_i) \times w_{\text{global}}(t_j) \\ &= \text{tf}(t_j, d_i) \times \log \left(\frac{N}{\text{df}(t_j)} \right) \end{aligned}$$

- In particular: weight zero for $\text{tf}(t_j, d_i)=0$ or $\text{df}(t_j)=N$
- Similarity: Cosine measure

$$\text{sim}(d_i, q) = \cos(\vec{d}_i, \vec{q}) = \frac{\vec{d}_i \cdot \vec{q}}{|\vec{d}_i| \cdot |\vec{q}|}$$

- Includes length normalization

Example

- Global weights

- E.g. for „cup“

$$w_{\text{global}}(\text{cup}) = \log\left(\frac{N}{df(\text{cup})}\right)$$

$$\log\left(\frac{5}{3}\right) \approx 0.22$$

1. coffee, coffee
2. cup, jar, jar, tea, tea
3. coffee, cup, cup, jar
4. coffee, coffee, coffee, cup, cup, cup, jar, jar, jar, tea
5. jar, jar, water, water

t_j	$df(t_j)$	$w_{\text{global}}(t_j)$
coffee	3	0.22
cup	3	0.22
jar	4	0.10
tea	2	0.40
water	1	0.70

Example

- Combined with local weights

t_j	d_1	d_2	d_3	d_4	d_5
coffee	2		1	3	
cup		1	2	3	
jar		2	1	3	2
tea		2		1	
water					2

$$\vec{d}_2 = (0 \cdot 0.22, 1 \cdot 0.22, 2 \cdot 0.10, 2 \cdot 0.4, 0 \cdot 0.7)$$

t_j	d_1	d_2	d_3	d_4	d_5
coffee	0.44	0	0.22	0.66	0
cup	0	0.22	0.44	0.66	0
jar	0	0.20	0.10	0.30	0.20
tea	0	0.80	0	0.40	0
water	0	0	0	0	1.40

- coffee, coffee
- cup, jar, jar, tea, tea
- coffee, cup, cup, jar
- coffee, coffee, coffee, cup,
cup, cup, jar, jar, tea
- jar, jar, water, water

t_j	$df(t_j)$	$w_{\text{global}}(t_j)$
coffee	3	0.22
cup	3	0.22
jar	4	0.10
tea	2	0.40
water	1	0.70

Query vector

- Construction of query vector
 - Just like a document
 - Using global weights from corpus

$$\begin{aligned} w_{\text{TF.IDF}}(t_j, q) &= w_{\text{local}}(t_j, q) \times w_{\text{global}}(t_j) \\ &= \text{tf}(t_j, q) \times \log \left(\frac{N}{\text{df}(t_j)} \right) \end{aligned}$$

- Note
 - Frequency of query terms matters
 - Sequence does not matter (bag of words)

Example

- Query: „cup jar“

- Query vector

$$\vec{q} = (0, 0.22, 0.1, 0, 0)$$

- Vector lengths:

$$|\vec{q}| = \sqrt{0.22^2 + 0.1^2} = 0.24$$

$$|\vec{d}_1| = 0.44$$

$$|\vec{d}_2| = 0.85$$

$$|\vec{d}_3| = 0.50$$

$$|\vec{d}_4| = 1.06$$

$$|\vec{d}_5| = 1.41$$

1. coffee, coffee
2. cup, jar, jar, tea, tea
3. coffee, cup, cup, jar
4. coffee, coffee, coffee, cup,
cup, cup, jar, jar, jar, tea
5. jar, jar, water, water

t_j	d_1	d_2	d_3	d_4	d_5
coffee	0.44	0	0.22	0.66	0
cup	0	0.22	0.44	0.66	0
jar	0	0.20	0.10	0.30	0.20
tea	0	0.80	0	0.40	0
water	0	0	0	0	1.40

Example

- Query vector:

$$\vec{q} = (0, 0.22, 0.1, 0, 0)$$

- Relevance (e.g. d_5):

$$\rho(\vec{d}_5, \vec{q}) = \frac{0 \cdot 0 + 0.22 \cdot 0 + 0.1 \cdot 0.2 + 0 \cdot 0 + 0 \cdot 1.4}{0.24 \cdot 1.41}$$

- Ranking

Rank	Document	ρ
1	d_3	0.89
2	d_4	0.69
3	d_2	0.33
4	d_5	0.05

- coffee, coffee
- cup, jar, jar, tea, tea
- coffee, cup, cup, jar
- coffee, coffee, coffee, cup, cup, cup, jar, jar, tea
- jar, jar, water, water

t_j	d_1	d_2	d_3	d_4	d_5
coffee	0.44	0	0.22	0.66	0
cup	0	0.22	0.44	0.66	0
jar	0	0.20	0.10	0.30	0.20
tea	0	0.80	0	0.40	0
water	0	0	0	0	1.40
Length	0.44	0.85	0.50	1.06	1.41


d_1 not in result list
as $\rho(d_1, q) = 0$

Summary – vector space ranking

- Represent the query as a weighted tf-idf vector
- Represent each document as a weighted tf-idf vector
- Compute the cosine similarity score for the query vector and each document vector
- Rank documents with respect to the query by score
- Return the top K (e.g., $K = 10$) to the user

› 3. Relevance Feedback (Rocchio)

Relevance Feedback for VSM



Incorporate Feedback

- User provides for some documents if they are relevant or irrelevant:
 - D^p : set of documents with positive feedback (relevant)
 - D^n : set of documents with negative feedback (irrelevant)
- Adjust query vector

$$\vec{q}_{FB} = \alpha \cdot \vec{q} + \beta \frac{1}{|D^p|} \sum_{d_i \in D^p} \vec{d}_i - \gamma \frac{1}{|D^n|} \sum_{d_i \in D^n} \vec{d}_i$$

- Parameters: typically $\alpha > \beta > \gamma$
- Adjust negative term weights to 0

Example

- Query:
 - paris hilton
- Documents

	paris	hilton	hotel	france	eiffel	blonde	heiress	actress
d ₁	3	1	1	2	1			
d ₂	1	3	4	1	3			
d ₃	2	1		1				
d ₄	3			2			1	
d ₅		1	3					
<hr/>								
d ₆	3	3				2		
d ₇	2	2					2	
d ₈	2	1	1			1	1	
d ₉	3	2				1		4
d ₁₀	3	2		1			2	3

Example

- Initial result list based on VSM
- User provides relevance feedback
 - Positive 
 - Negative 
- Parameter:
 - $\alpha = 1$
 - $\beta = 0.75$
 - $\gamma = 0.15$

Rank	Doc	ρ
1	d_3	0.395
2	d_6	0.183
3	d_7	0.161
4	d_4	0.132
5	d_1	0.128
6	d_8	0.125
7	d_{10}	0.071
8	d_9	0.057
9	d_2	0.049
10	d_5	0.027



Example

- Adjust query

$$\begin{aligned} \alpha &= 1 \\ \beta &= 0.75 \\ \gamma &= 0.15 \end{aligned}$$

$$\vec{q}_{FB} = 1 \cdot \begin{pmatrix} 0.046 \\ 0.046 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \end{pmatrix} + 0.75 \cdot \frac{1}{2} \cdot \begin{pmatrix} 0.137 \\ 0.046 \\ 0.398 \\ 0.602 \\ 0.699 \\ 0 \\ 0 \\ 0 \end{pmatrix} + \begin{pmatrix} 0.092 \\ 0.092 \\ 0 \\ 0.301 \\ 0 \\ 0 \\ 0 \\ 0 \end{pmatrix}$$

$$-0.15 \cdot \frac{1}{2} \cdot \begin{pmatrix} 0.092 \\ 0.092 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0.796 \\ 0 \end{pmatrix} + \begin{pmatrix} 0.092 \\ 0.046 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0.523 \\ 0.398 \end{pmatrix} \mapsto \begin{pmatrix} 0.118 \\ 0.087 \\ 0.119 \\ 0.339 \\ 0.262 \\ 0 \\ 0 \\ 0 \end{pmatrix}$$

Rank	Doc	ρ
1	d_1	0.957
2	d_3	0.787
3	d_2	0.691
4	d_4	0.640
5	d_5	0.262
6	d_8	0.172
7	d_{10}	0.119
8	d_6	0.056
9	d_7	0.050
10	d_9	0.018

- Retrieve new results

Relevance Feedback in Practice

- Users unwilling to give feedback
 - Query reformulation is easier
- Pseudo relevance feedback
 - Retrieve result list (do not show to user)
 - Use top-k results as positive feedback
 - Rarely low-ranking documents as negative feedback
 - Adjust query vector
 - Retrieve final result list
- Works good, when initial results are good

> 4. Summary

Summary

- At the end of this lecture you should understand the following concepts
 - Boolean Model
 - Ranked retrieval model
 - Scoring
 - Term frequency
 - Document frequency
 - TF-IDF
 - Vector Space Model
 - Relevance feedback

References

- [1] <https://olat.vcrp.de/auth/RepositoryEntry/4071063853>
- [2] <https://nlp.stanford.edu/IR-book/information-retrieval-book.html> Christopher D. Manning, Prabhakar Raghavan and Hinrich Schütze, Introduction to Information Retrieval, Cambridge University Press. 2008.
 - ▶ Chapter 5 (VSM) and Chapter 9 (Relevance Feedback)