

# ➤ Web Information Retrieval

## Web Crawling

### (SOSE 2023, 07.06.2023)

Frank Hopfgartner, Stefania Zourlidou  
Institute for Web Science and Technologies

# Credit for these slides

These slides have been adapted from

- Web IR (Zeyd Boukhers-WeST, SOSE 2020)

# Recapitulation

- Classical IR vs. Web IR
- Web IR, Web search basics
  - Ads
  - Spams
  - Duplicates
  - User Needs
  - Web Graph (anchor text)
  - Indexing

# Objectives of this lecture


- Crawler
  - what is it?
    - features a crawler *must* provide
    - features a crawler *should* provide
  - crawler architecture
    - robots exclusion protocol
    - url normalization
  - why distributing the crawler
  - the URL frontier

# ➤ 1. Web crawling: what is it?

# Basic crawler operation

1. Begin with a known “seed” URLs
2. Fetch and parse the URLs
  - a) Extract URLs they point to
  - b) Place the extracted URLs on the queue
3. Go to 2

# Crawling picture



# Simple picture – complications

- Web crawling is not feasible with one machine
  - All of the above steps are distributed
- Malicious pages
  - Spam pages
  - Spider traps – incl dynamically generated
- Even non-malicious pages pose challenges
  - Latency/bandwidth to remote servers vary
  - Webmasters' stipulations
    - How “deep” should you crawl a site's URL hierarchy?
  - Site mirrors and duplicate pages
- Politeness – don't hit a server too often


# What a crawler **must** do

- Be **polite**: respect implicit and explicit politeness considerations
  - Only crawl allowed pages
  - Respect *robots.txt*
- Be **robust**: be immune to spider traps and other malicious behavior from web servers

# What a crawler **should** do

- Be capable of **distributed** operation: designed to run on multiple distributed machines
- Be **scalable**: designed to increase the crawl rate by adding more machines
- **Performance/efficiency**: permit full use of available processing and network resources
- Fetch pages of “higher **quality**” first
- **Continuous** operation: Continue fetching fresh copies of a previously fetched page
- **Extensible**: Adapt to new data formats, protocols

# Updated crawling picture



- Can include multiple pages from the same host
- Must avoid trying to fetch them all at the same time
- Must try to keep all crawling threads busy

# Explicit and implicit politeness

- **Explicit politeness:** specifications from webmasters on what portions of site can be crawled
  - robots.txt
- **Implicit politeness:** even with no specification, avoid hitting any site too often

- Protocol for giving spiders (“robots”) limited access to a website, originally from 1994
  - [www.robotstxt.org/wc/norobots.html](http://www.robotstxt.org/wc/norobots.html)
- Website announces its request on what can(not) be crawled
  - For a server, create a file /robots.txt
  - This file specifies access restrictions

# Robots.txt example

- No robot should visit any URL starting with "/yoursite/temp/", except the robot called "searchengine"

User-agent: \*

Disallow: /yoursite/temp/

User-agent: searchengine

Disallow:

# Processing steps in crawling


- Pick a URL from the frontier
- Fetch the document (web page) at the URL
- Parse the URL
  - Extract the text and set of links from the fetched web page
- Check if URL has content already seen
  - If not, add to indexes
- For each extracted URL
  - Ensure it passes certain URL filter tests
  - Check if it is already in the frontier (duplicate URL elimination)

Which one?

E.g., only crawl .edu,  
obey robots.txt, etc.

## ➤ 2. Basic crawl architecture

# Basic crawl architecture



- A lookup service on the internet
  - Given a URL, retrieve its IP address
  - Service provided by a distributed set of servers – thus, lookup latencies can be high (even seconds)
- Common OS implementations of DNS lookup are *blocking*: only one outstanding request at a time
- Solutions
  - DNS caching
  - Batch DNS resolver – collects requests and sends them out together

# Parsing: URL normalization

- When a fetched document is parsed, some of the extracted links are *relative* URLs
  - E.g., [http://en.wikipedia.org/wiki/Main\\_Page](http://en.wikipedia.org/wiki/Main_Page) has a relative link to /wiki/Wikipedia:General\_disclaimer which is the same as the absolute URL  
[http://en.wikipedia.org/wiki/Wikipedia:General\\_disclaimer](http://en.wikipedia.org/wiki/Wikipedia:General_disclaimer)
  - During parsing, must normalize (expand) such relative URLs

# Content seen?

- Duplication is widespread on the web
- If the page just being fetched is already in the index, do not further process it
- This is verified using document fingerprints or shingles

- Filters – regular expressions for URL's to be crawled/not
- Once a `robots.txt` file is fetched from a site, need not fetch it repeatedly
  - Doing so burns bandwidth, hits web server
- Cache `robots.txt` files

# Duplicate URL elimination

- For a non-continuous (one-shot) crawl, test to see if an extracted+filtered URL has already been passed to the frontier
- For a continuous crawl – see details of frontier implementation

## › 3. Distributing the crawler

# Distributing the crawler


- Run multiple crawl threads, under different processes – potentially at different nodes
  - Geographically distributed nodes
- Partition hosts being crawled into nodes
  - Hash used for partition
- How do these nodes communicate and share URLs?

# Google data centers (wazfaring. com)



# Communication between nodes

Output of the URL filter at each node is sent to the Dup URL Eliminator of the appropriate node



## > 4. The URL frontier

# URL frontier: two main considerations

- **Politeness:** do not hit a web server too frequently
- **Freshness:** crawl some pages more often than others
  - E.g., pages (such as News sites) whose content changes often


These goals may conflict each other

- E.g., simple priority queue fails – many links out of a page go to its own site, creating a burst of accesses to that site


# Politeness – challenges

- Even if we restrict only one thread to fetch from a host, can hit it repeatedly
- Common heuristic: insert time gap between successive requests to a host that is  $\gg$  time for most recent fetch from that host

# URL frontier: Mercator scheme




# Mercator URL frontier



- URLs flow in from the top into the frontier
- Front queues manage prioritization
- Back queues enforce politeness
- Each queue is FIFO

# Front queues




- Prioritizer assigns to URL an integer priority between 1 and  $K$ 
  - Appends URL to corresponding queue
- Heuristics for assigning priority
  - Refresh rate sampled from previous crawls
  - Application-specific (e.g., “crawl news sites more often”)

# Biased front queue selector

- When a back queue requests a URL (in a sequence to be described): picks a **front queue** from which to pull a URL
- This choice can be round robin biased to queues of higher priority, or some more sophisticated variant
  - Can be randomized

# Back queues



# Back queue invariants

- Each back queue is kept non-empty while the crawl is in progress
- Each back queue only contains URLs from a single host
  - Maintain a table from hosts to back queues

Host name	Back queue
...	3
	1
	$B$

- One entry for each back queue
- The entry is the earliest time  $t_e$  at which the host corresponding to the back queue can be hit again
- This earliest time is determined from
  - Last access to that host
  - Any time buffer heuristic we choose

- A crawler thread seeking a URL to crawl
  - Extracts the root of the heap
  - Fetches URL at head of corresponding back queue  $q$  (look up from table)
  - Checks if queue  $q$  is now empty – if so, pulls a URL  $v$  from front queues
    - If there is already a back queue for  $v$ 's host, append  $v$  to  $q$  and pull another URL from front queues, repeat
    - Else add  $v$  to  $q$
  - When  $q$  is non-empty, create heap entry for it

# Number of back queues $B$

- Keep all threads busy while respecting politeness
- Mercator recommendation: three times as many back queues as crawler threads

## › 5. Summary

- Web crawling
  - what is it?
  - basic crawler architecture
  - distributing the basic crawler architecture
  - the URL frontier

- [1] <https://olat.vcrp.de/auth/RepositoryEntry/4071063853>
- [2] <https://nlp.stanford.edu/IR-book/information-retrieval-book.html>  
Christopher D. Manning, Prabhakar Raghavan and Hinrich Schütze,  
Introduction to Information Retrieval, Cambridge University Press. 2008
  - ▶ Chapter 20 (Web crawling and indexes)
- [3] Some extra resources
  - ▶ [Mercator: A scalable, extensible Web crawler](#)
  - ▶ [A Standard for Robot Exclusion](#)